

# Functional PCA in R

1 October 2021

## Overview

This is a brief introduction to the package `fdapace` (Gajardo et al. 2021). For a general overview on functional data analysis (FDA) see (Wang, Chiou, and Müller 2016) and key references for the PACE approach and the associated dynamics are (Yao et al. 2003; Yao, Müller, and Wang 2005; Liu and Müller 2009; Müller and Yao 2010; Li and Hsing 2010; Zhang and Wang 2016; Zhang and Wang 2018). The basic work-flow behind the PACE approach for sparse functional data is as follows (see e.g. (Yao, Müller, and Wang 2005; Liu and Müller 2009) for more information):

1. Calculate the smoothed mean  $\hat{\mu}$  (using local linear smoothing) aggregating all the available readings together.
2. Calculate for each curve separately its own raw covariance and then aggregate all these raw covariances to generate the sample raw covariance.
3. Use the off-diagonal elements of the sample raw covariance to estimate the smooth covariance.
4. Perform eigenanalysis on the smoothed covariance to obtain the estimated eigenfunctions  $\hat{\phi}$  and eigenvalues  $\hat{\lambda}$ , then project that smoothed covariance on a positive semi-definite surface (Hall, Müller, and Yao 2008).
5. Use Conditional Expectation (PACE step) to estimate the corresponding scores  $\hat{\xi}$ . ie.

$$\hat{\xi}_{ik} = \hat{E}[\hat{\xi}_{ik}|Y_i] = \hat{\lambda}_k \hat{\phi}_{ik}^T \Sigma_{Y_i}^{-1} (Y_i - \hat{\mu}_i).$$

As a working assumption a dataset is treated as sparse if it has on average less than 20, potentially irregularly sampled, measurements per subject. A user can manually change the automatically determined `dataType` if that is necessary. For densely observed functional data simplified procedures are available to obtain the eigencomponents and associated functional principal components scores (see eg. (Castro, Lawton, and Sylvestre 1986) for more information). In particular in this case we:

1. Calculate the cross-sectional mean  $\hat{\mu}$ .
2. Calculate the cross-sectional covariance surface (which is guaranteed to be positive semi-definite).
3. Perform eigenanalysis on the covariance to estimate the eigenfunctions  $\hat{\phi}$  and eigenvalues  $\hat{\lambda}$ .
4. Use numerical integration to estimate the corresponding scores  $\hat{\xi}$ . ie.  $\hat{\xi}_{ik} = \int_0^T [y(t) - \hat{\mu}(t)] \phi_i(t) dt$

In the case of sparse FPCA the most computational intensive part is the smoothing of the sample's raw covariance function. For this, we employ a local weighted bilinear smoother.

A sibling MATLAB package for `fdapace` can be found [here](#).

## FPCA in R using `fdapace`

The simplest scenario is that one has two lists `yList` and `tList` where `yList` is a list of vectors, each containing the observed values  $Y_{ij}$  for the  $i$ th subject and `tList` is a list of vectors containing corresponding time points. In this case one uses:

```
FPCAobj <- FPCA(Ly=yList, Lt=tList)
```

The generated FPCAobj will contain all the basic information regarding the desired FPCA.

## Generating a toy dense functional dataset from scratch

```
library(fdapace)

# Set the number of subjects (N) and the
# number of measurements per subjects (M)
N <- 200;
M <- 100;
set.seed(123)

# Define the continuum
s <- seq(0,10,length.out = M)

# Define the mean and 2 eigenvectors
meanFunct <- function(s) s + 10*exp(-(s-5)^2)
eigFunct1 <- function(s) +cos(2*s*pi/10) / sqrt(5)
eigFunct2 <- function(s) -sin(2*s*pi/10) / sqrt(5)

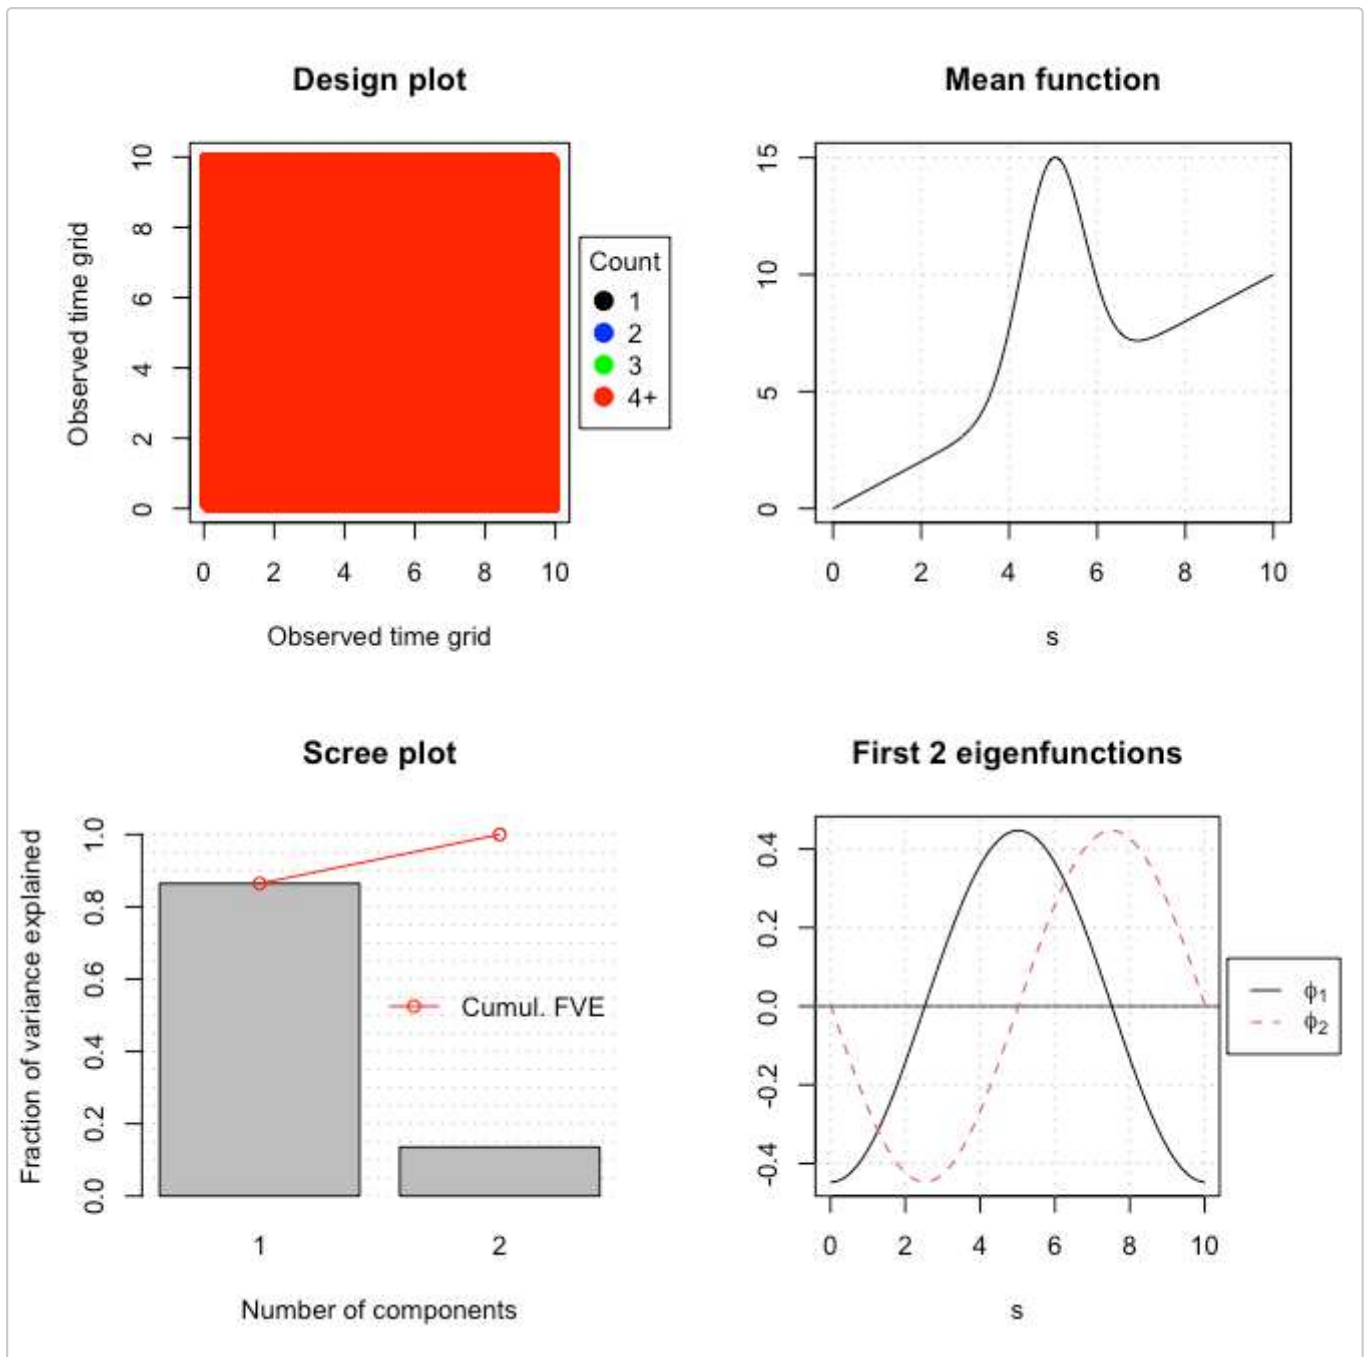
# Create FPC scores
Ksi <- matrix(rnorm(N*M), ncol=2);
Ksi <- apply(Ksi, 2, scale)
Ksi <- Ksi %*% diag(c(5,2))

# Create Y_true
yTrue <- Ksi %*% t(matrix(c(eigFunct1(s),eigFunct2(s)), ncol=2)) +
t(matrix(rep(meanFunct(s),N), nrow=M))
```

## Running FPCA on a dense dataset

```
L3 <- MakeFPCAInputs(IDs = rep(1:N, each=M), tVec=rep(s,N), t(yTrue))
FPCAdense <- FPCA(L3$Ly, L3$Lt)

# Plot the FPCA object
plot(FPCAdense)
```



```
# Find the standard deviation associated with each component
sqrt(FPCAdense$lambda)
```

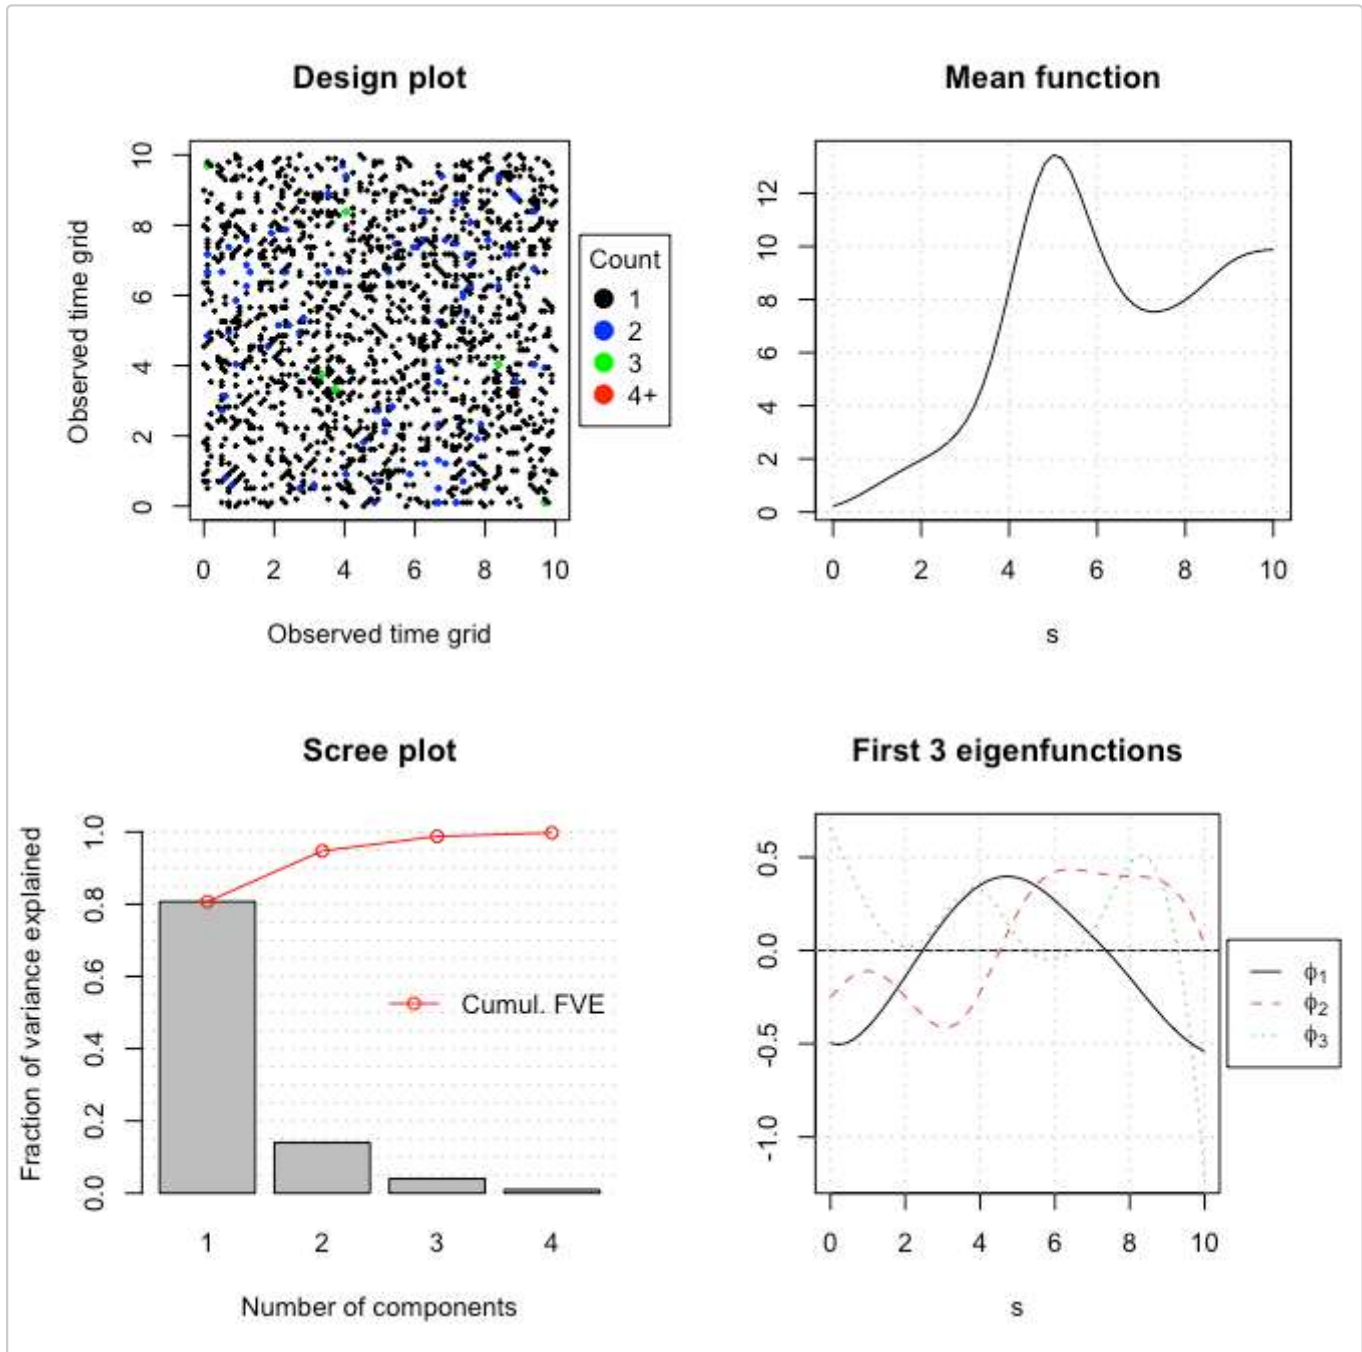
```
## [1] 5.050606 1.999073
```

## Running FPCA on a sparse and noisy dataset

```
# Create sparse sample
# Each subject has one to five readings (median: 3)
set.seed(123)
ySparse <- Sparsify(yTrue, s, sparsity = c(1:5))
```

```
# Give your sample a bit of noise
ySparse$yNoisy <- lapply(ySparse$Ly, function(x) x + 0.5*rnorm(length(x)))

# Do FPCA on this sparse sample
# Notice that sparse FPCA will smooth the data internally (Yao et al., 2005)
# Smoothing is the main computational cost behind sparse FPCA
FPCAsparse <- FPCA(ySparse$yNoisy, ySparse$Lt, list(plot = TRUE))
```



## Further functionality

FPCA calculates the bandwidth utilized by each smoother using generalised cross-validation or  $k$ -fold cross-validation automatically. Dense data are not smoothed by default. The argument `methodMuCovEst` can be

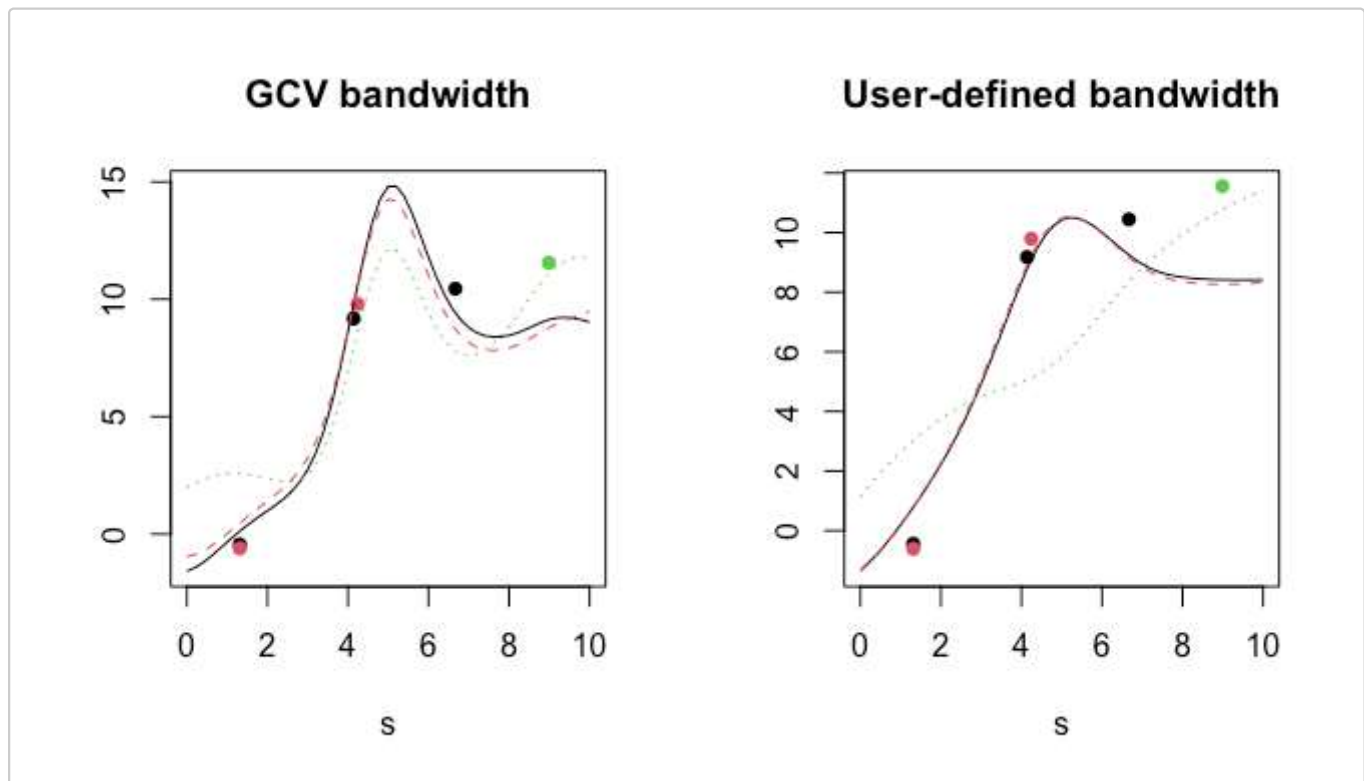
switched between `smooth` and `cross-sectional` if one wants to utilize different estimation techniques when work with dense data.

The bandwidth used for estimating the smoothed mean and the smoothed covariance are available under `...bwMu` and `bwCov` respectively. Users can nevertheless provide their own bandwidth estimates:

```
FPCAsparseMuBW5 <- FPCA(ySparse$yNoisy, ySparse$Lt, optns= list(userBwMu = 5))
```

Visualising the fitted trajectories is a good way to see if the new bandwidth made any sense:

```
par(mfrow=c(1,2))
CreatePathPlot( FPCAsparse, subset = 1:3, main = "GCV bandwidth", pch = 16)
CreatePathPlot( FPCAsparseMuBW5, subset = 1:3, main = "User-defined bandwidth", pch = 16)
```



FPCA uses a Gaussian kernel when smoothing sparse functional data; other kernel types (eg. Epanechnikov/`epan`) are also available (see `?FPCA`). The kernel used for smoothing the mean and covariance surface is the same. It can be found under `$optns$kernel` of the returned object. For instance, one can switch the default Gaussian kernel (`gauss`) for a rectangular kernel (`rect`) as follows:

```
FPCAsparseRect <- FPCA(ySparse$yNoisy, ySparse$Lt, optns = list(kernel = 'rect')) # Use rectangular kernel
```

FPCA returns automatically the smallest number of components required to explain 99% of a sample's variance. Using the function `selectK` one can determine the number of relevant components according to AIC, BIC or a different Fraction-of-Variance-Explained threshold. For example:

```
SelectK( FPCAsparse, criterion = 'FVE', FVEthreshold = 0.95) # K = 2
```

```
## $K
## [1] 3
##
## $criterion
## [1] 0.9876603
```

```
SelectK( FPCAsparse, criterion = 'AIC') # K = 2
```

```
## $K
## [1] 2
##
## $criterion
## [1] 993.4442
```

When working with functional data (usually not very sparse) the estimation of derivatives is often of interest. Using `fitted.FPCA` one can directly obtain numerical derivatives by defining the appropriate order `p`; `fdapace` provides for the first two derivatives (`p = 1` or `2`). Because the numerically differentiated data are smoothed the user can define smoothing specific arguments (see `?fitted.FPCA` for more information); the derivation is done by using the derivative of the linear fit. Similarly using the function `FPCAder`, one can augment an `FPCA` object with functional derivatives of a sample's mean function and eigenfunctions.

```
fittedCurvesP0 <- fitted(FPCAsparse) # equivalent: fitted(FPCAsparse, derOpts=list(p = 0));
# Get first order derivatives of fitted curves, smooth using Epanechnikov kernel
fittedCurvesP1 <- fitted(FPCAsparse, derOpts=list(p = 1, kernelType = 'epan'))
```

```
## Warning in fitted.FPCA(FPCAsparse, derOpts = list(p = 1, kernelType = "epan")): Potentially you
use too many components to estimate derivatives.
## Consider using SelectK() to find a more informed estimate for 'K'.
```

## A real-world example

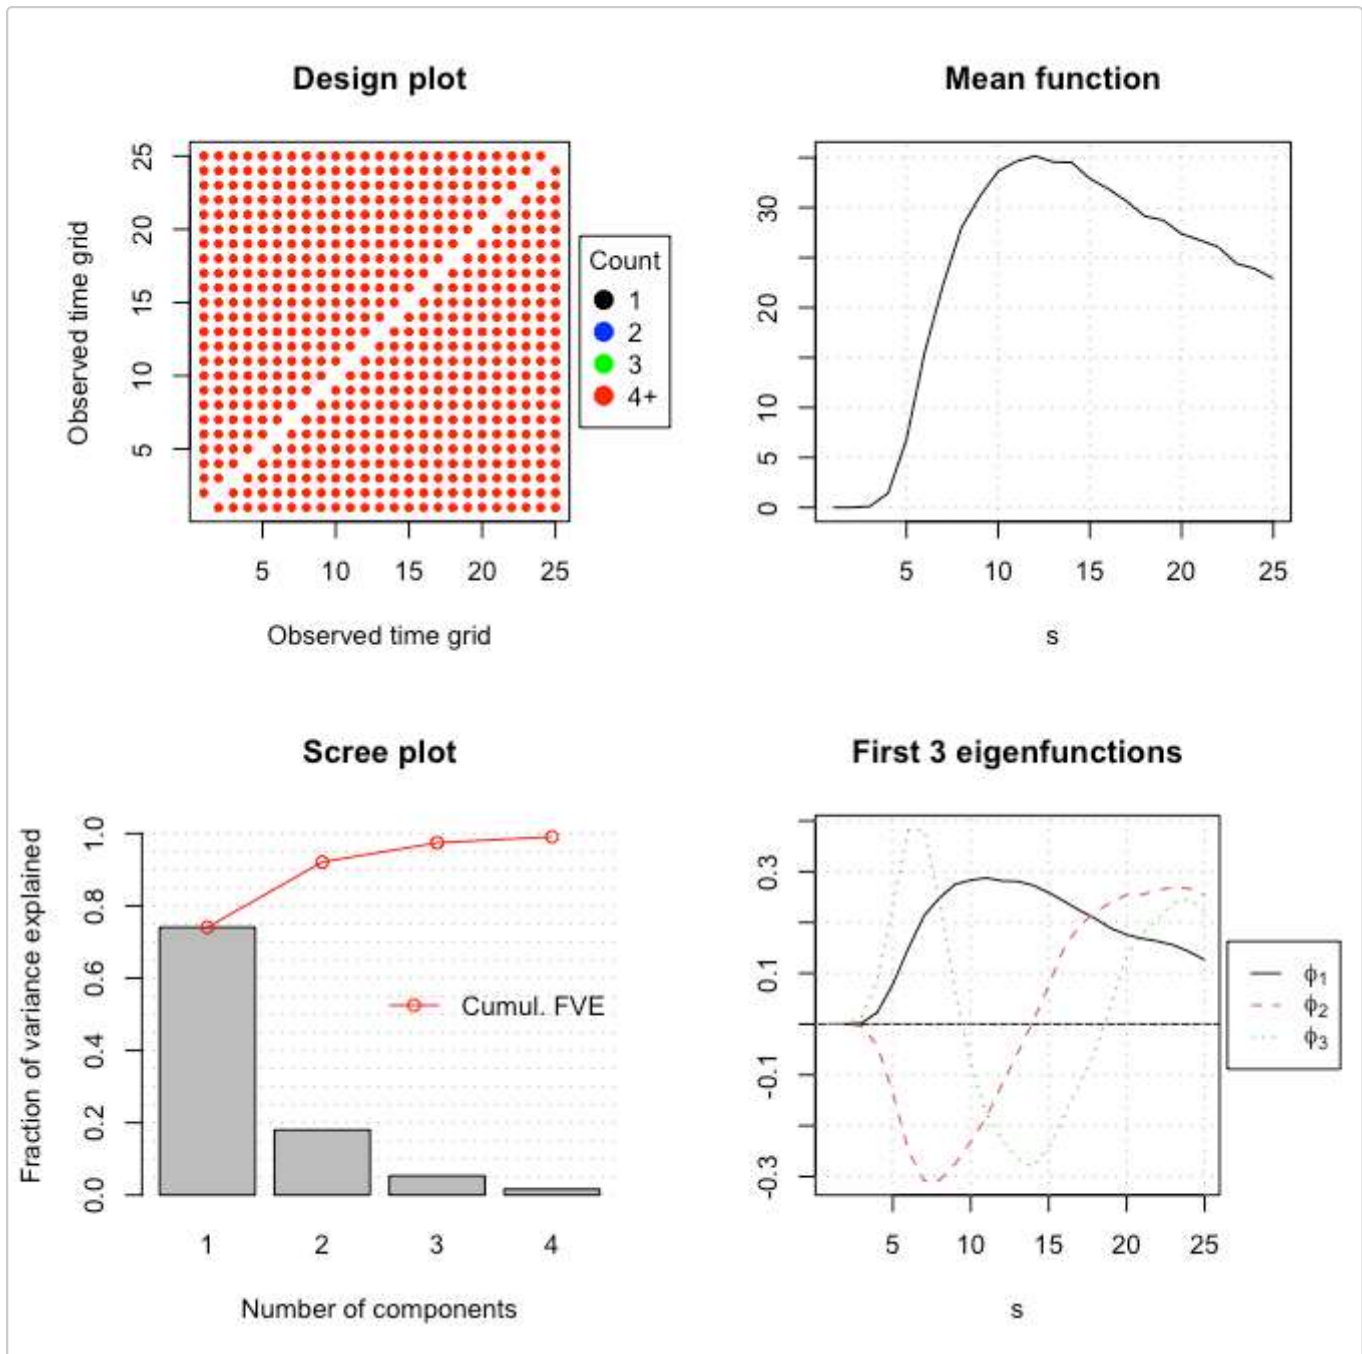
We use the `medfly25` dataset that is available with `fdapace` to showcase `FPCA` and its related functionality. `medfly25` is a dataset containing the eggs laid from 789 medflies (Mediterranean fruit flies, *Ceratitis capitata*) during the first 25 days of their lives. It is a subset of the dataset used by Carey et al. (1998) (Carey et al. 1998); only flies having lived at least 25 days are shown. The data are rather noisy, dense and with a characteristic flat start. For that reason in contrast with above we will use a smoothing estimating procedure despite having dense data.

```
# Load data
data(medfly25)

# Turn the original data into a list of paired amplitude and timing lists
Flies <- MakeFPCAInputs(medfly25$ID, medfly25$Days, medfly25$nEggs)
```



```
fPCAObjFlies <- FPCA(Flies$Ly, Flies$Lt, list(plot = TRUE, methodMuCovEst = 'smooth',
userBwCov = 2))
```



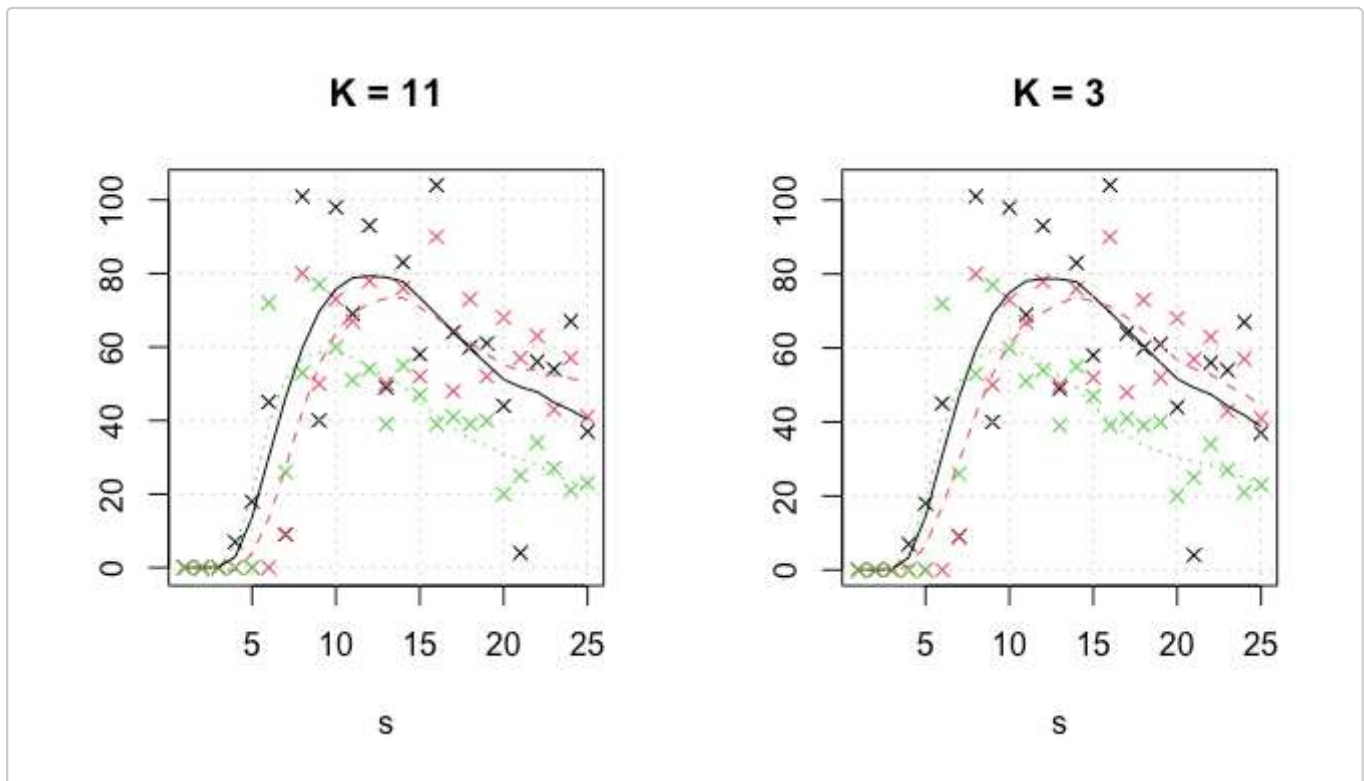
Based on the scree-plot we see that the first three components appear to encapsulate most of the relevant variation. The number of eigencomponents to reach a 99.99% FVE is 11 but just 3 eigencomponents are enough to reach a 95.0%. We can easily inspect the following visually, using the `CreatePathPlot` command.

```
require('ks')
```

```
## Loading required package: ks
```

```
par(mfrow=c(1,2))
CreatePathPlot(fPCAObjFlies, subset = c(3,5,135), main = 'K = 11', pch = 4); grid()
```

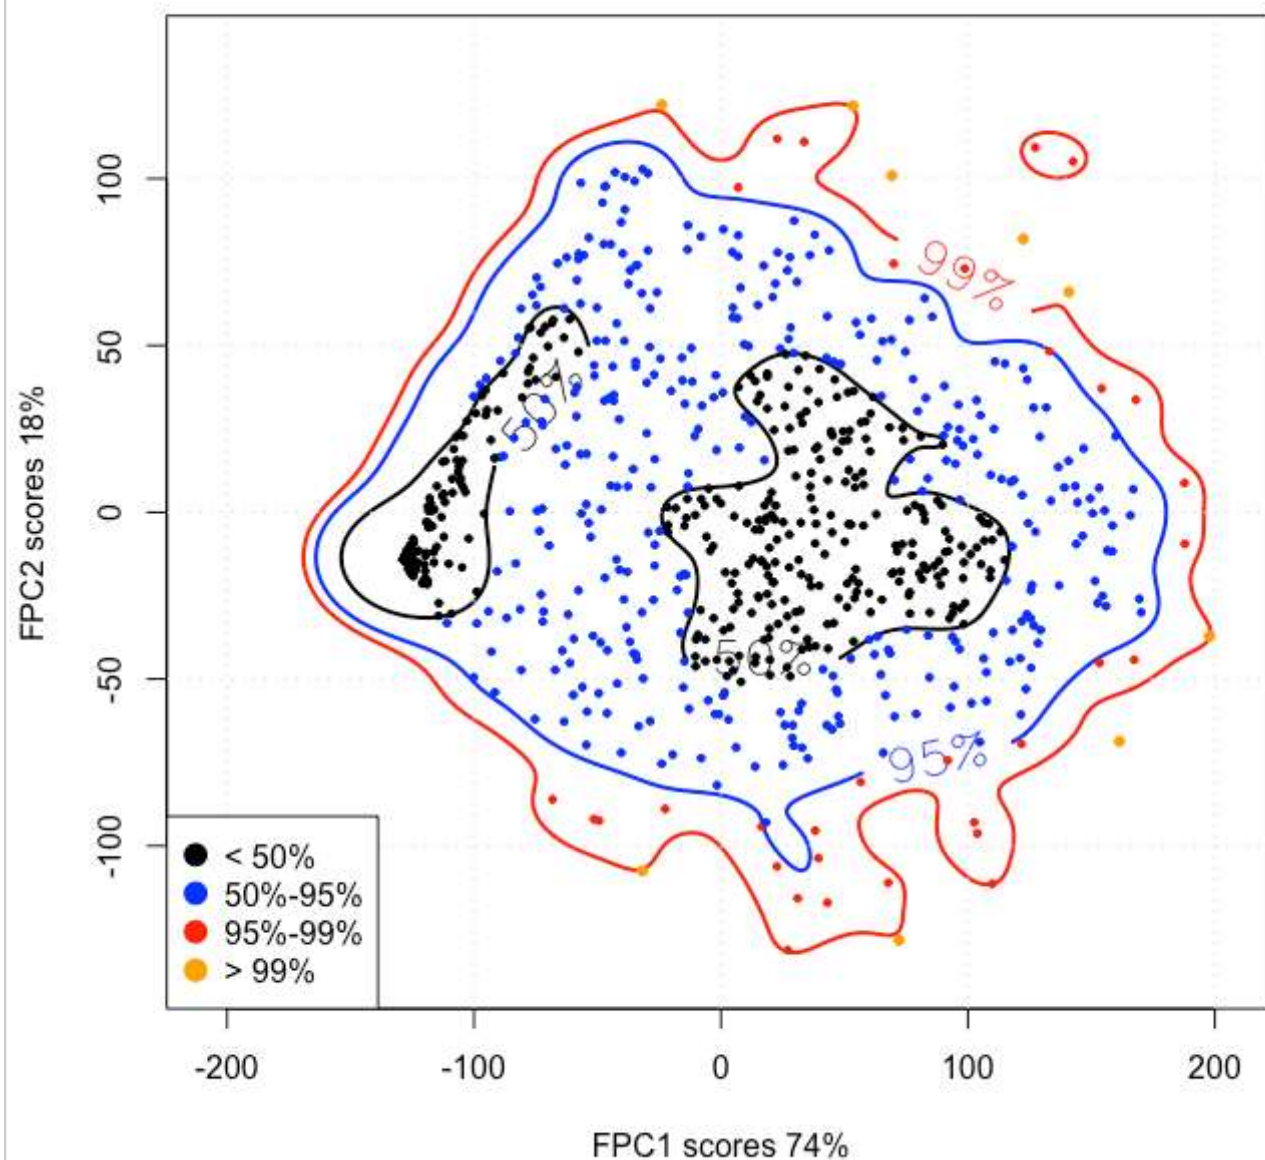
```
CreatePathPlot(fpcaObjFlies, subset = c(3,5,135), K = 3, main = 'K = 3', pch = 4) ; grid()
```




One can perform outlier detection (Febrero, Galeano, and González-Manteiga 2007) as well as visualize data using a functional box-plot. To achieve these tasks one can use the functions `CreateOutliersPlot` and `CreateFuncBoxPlot`. Different ranking methodologies (KDE, bagplot Hyndman and Shang (2010) or point-wise) are available and can potentially identify different aspects of a sample. For example here it is notable that the kernel density estimator KDE variant identifies two main clusters within the main body of sample. By construction the bagplot method would use a single bag and this feature would be lost. Both functions return a (temporarily) invisible copy of a list containing the labels associated with each of sample curve0. `CreateOutliersPlot` returns a (temporarily) invisible copy of a list containing the labels associated with each of sample curve.

```
par(mfrow=c(1,1))
CreateOutliersPlot(fpcaObjFlies, optns = list(K = 3, variant = 'KDE'))
```



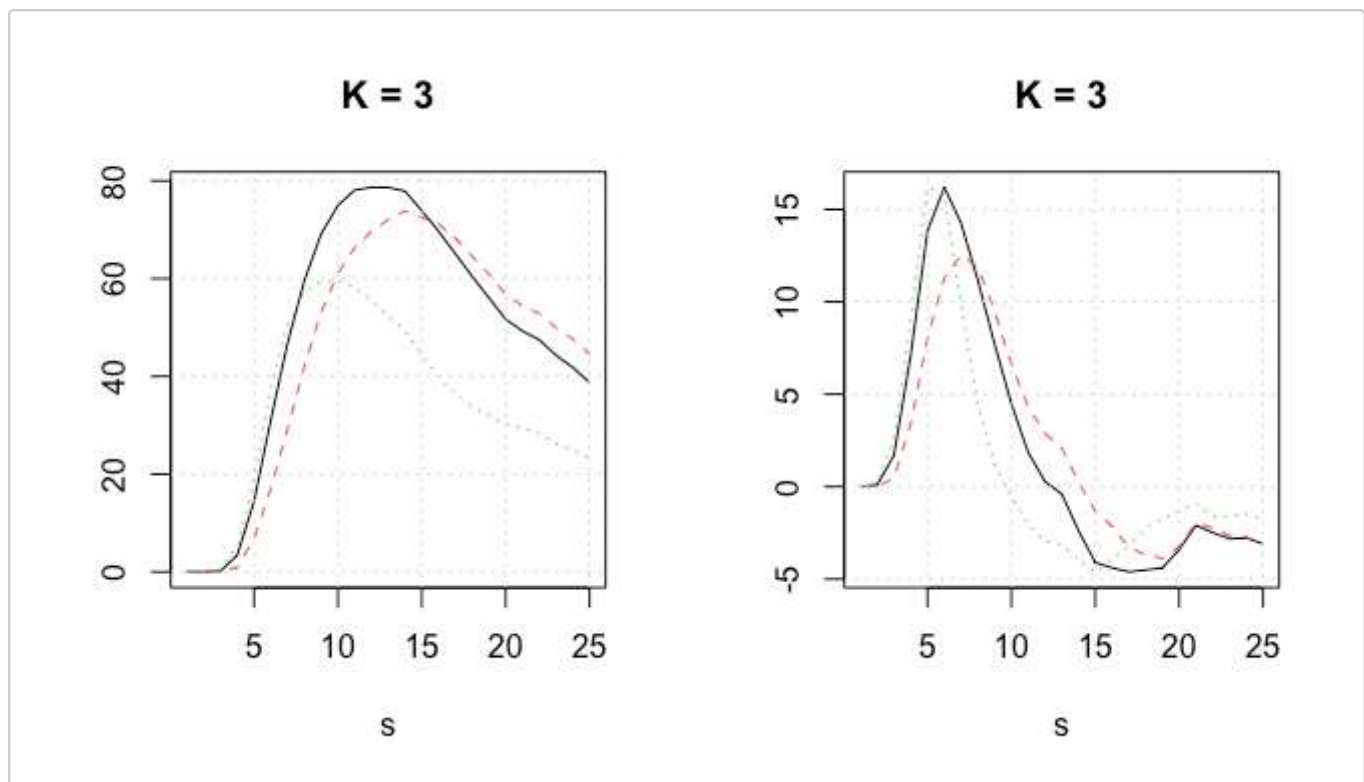


```
CreateFuncBoxPlot(fpcaObjFlies, xlab = 'Days', ylab = '# of eggs laid', optns = list(K = 3,
variant='bagplot'))
```



Functional data lend themselves naturally to questions about their rate of change; their derivatives. As mentioned previously using `fdapace` one can generate estimates of the sample's derivatives ( `fitted.FPCA`) or the derivatives of the principal modes of variation ( `FPCAder`). In all cases, one defines a `derOptsns` list of options to control the derivation parameters. Getting derivatives is obtained by using a local linear smoother as above.

```
par(mfrow=c(1,2))
CreatePathPlot(fpcaObjFlies, subset = c(3,5,135), K = 3, main = 'K = 3', showObs = FALSE)
; grid()
CreatePathPlot(fpcaObjFlies, subset = c(3,5,135), K = 3, main = 'K = 3', showObs = FALSE,
derOptsns = list(p = 1, bw = 1.01 , kernelType = 'epan') ) ; grid()
```



## References

- Carey, JR, P Liedo, H-G Müller, J-L Wang, and J-M Chiou. 1998. "Relationship of Age Patterns of Fecundity to Mortality, Longevity, and Lifetime Reproduction in a Large Cohort of Mediterranean Fruit Fly Females." *The Journals of Gerontology Series A: Biological Sciences and Medical Sciences* 53 (4): B245–51.
- Castro, PE, WH Lawton, and EA Sylvestre. 1986. "Principal Modes of Variation for Processes with Continuous Sample Curves." *Technometrics* 28 (4): 329–37.
- Febrero, M, P Galeano, and W González-Manteiga. 2007. "A Functional Analysis of NOx Levels: Location and Scale Estimation and Outlier Detection." *Computational Statistics* 22 (3): 411–27.
- Gajardo, Alvaro, Satarupa Bhattacharjee, Cody Carroll, Yaqing Chen, Xiongtao Dai, Jianing Fan, Pantelis Z. Hadjipantelis, et al. 2021. *fdapace: Functional Data Analysis and Empirical Dynamics*. <https://CRAN.R-project.org/package=fdapace>.
- Hall, P, H-G Müller, and F Yao. 2008. "Modelling Sparse Generalized Longitudinal Observations with Latent Gaussian Processes." *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 70 (4): 703–23.
- Hyndman, RJ, and HL Shang. 2010. "Rainbow Plots, Bagplots, and Boxplots for Functional Data." *Journal of Computational and Graphical Statistics* 19 (1).
- Li, Y, and T Hsing. 2010. "Uniform Convergence Rates for Nonparametric Regression and Principal Component Analysis in Functional/Longitudinal Data." *The Annals of Statistics* 38 (6): 3321–51.
- Liu, B, and H-G Müller. 2009. "Estimating Derivatives for Samples of Sparsely Observed Functions, with Application to Online Auction Dynamics." *Journal of the American Statistical Association* 104 (486): 704–17.
- Müller, H-G, and F Yao. 2010. "Empirical Dynamics for Longitudinal Data." *The Annals of Statistics* 38 (6): 3458–86.
- Rousseeuw, PJ, I Ruts, and JW Tukey. 1999. "The Bagplot: A Bivariate Boxplot." *The American Statistician* 53 (4): 382–87.
- Wang, J-L, J-M Chiou, and H-G Müller. 2016. "Functional Data Analysis." *Annual Review of Statistics and Its Application* 3: 257–95.

- Yao, F, H-G Müller, A J Clifford, S R Dueker, J Follett, Y Lin, B A Buchholz, and JS Vogel. 2003. "Shrinkage Estimation for Functional Principal Component Scores with Application to the Population Kinetics of Plasma Folate." *Biometrics* 59 (3): 676–85.
- Yao, F, H-G Müller, and J-L Wang. 2005. "Functional Data Analysis for Sparse Longitudinal Data." *Journal of the American Statistical Association* 100 (470): 577–90.
- Zhang, X, and J-L Wang. 2016. "From Sparse to Dense Functional Data and Beyond." *The Annals of Statistics* 44 (5): 2281–321.
- Zhang, X, and J-L Wang. 2018. "Optimal Weighting Schemes for Longitudinal and Functional Data." *Statistics & Probability Letters* 138: 165–70.